

Responsible AI Case Study

Fabricated legal citations from ChatGPT in *Mata v. Avianca, Inc.* (S.D.N.Y., 2023)

Incident category	Hallucination and misinformation
Deliverables	1. Case Study Report 2. Prevention Recommendations 3. Responsible AI Summary
Related specialization	Advanced RAG (grounding and citation verification)

1. Case Study Report

1.1 What happened

- Roberto Mata sued Avianca Airlines over an injury he said he suffered on a 2019 flight. His attorney, Steven Schwartz of Levidow, Levidow & Oberman, had been licensed in New York for over three decades.
- After Avianca moved to dismiss the case as untimely, Mata's lawyers filed an opposition brief citing multiple court decisions. Schwartz had used ChatGPT to find supporting case law.
- Avianca's lawyers told the court they could not locate most of the cited cases, and that the few they found did not support the claims made about them.
- The court could not find the cases either and ordered the lawyers to provide copies. Mata's lawyers supplied documents purporting to be the cases, after ChatGPT assured them the cases existed and could be found in databases such as LexisNexis and Westlaw.
- At least six decisions were bogus, with fake quotations and fake internal citations. Examples included *Varghese v. China South Airlines*, *Martinez v. Delta Airlines* and *Petersen v. Iran Air*.
- Schwartz stated in an affidavit that he had never used ChatGPT for legal research before and was unaware its content could be false.

1.2 Timeline

Date	Event
Early 2023	Opposition brief filed containing non-existent cases.
April 2023	Avianca's counsel tells the court it cannot locate the cited cases; the judge issues show-cause orders requiring copies.
May 2023	Judge Castel calls the situation an "unprecedented circumstance"; an order to show cause on sanctions follows.
June 8, 2023	Sanctions hearing.
June 22, 2023	Court imposes a \$5,000 penalty and dismisses the underlying suit as time-barred.

1.3 Root causes

- **Model behavior:** A large language model generates plausible-sounding text rather than looking facts up. It produced realistic case names, courts, judges and quotations that did not exist.
- **Misplaced trust in self-verification:** The lawyer asked the same model whether the cases were real. The model confidently said yes. A model cannot reliably audit its own output.
- **No human verification:** Nobody checked the citations against an authoritative legal database.

- **Poor incident response:** The judge noted the sanctions might not have been needed had the lawyers come clean promptly. They instead stood by the fake opinions after the court questioned them.

1.4 Impact

- **Legal and financial:** The firm and both lawyers were fined \$5,000 after the court found bad faith. They were also ordered to send the ruling to their client and to each judge falsely named as the author of a fake opinion. The judge declined to compel an apology.
- **Case outcome:** The suit was dismissed as time-barred, so the fabricated research did nothing for the client and consumed court time.
- **Reputational:** The case received wide press coverage and became the best-known example of AI hallucination in professional practice.
- **Systemic:** The judge wrote that using AI is not improper in itself, but lawyers have a gatekeeping duty to ensure the accuracy of filings. The case is widely credited with prompting courts and bar bodies to issue AI verification and disclosure rules.

2. Prevention Recommendations

Hallucination cannot be eliminated entirely, so prevention combines technical safeguards, process controls and organizational policy.

Level	Recommendation	Why it helps
Technical	Use retrieval-augmented generation (RAG) over a trusted source such as a legal database, so answers are drawn from real documents.	Grounds output in verifiable text instead of model memory.
	Add an automated citation verifier that checks every case name, court and citation against an authoritative index before output is shown.	Would have caught all six fake cases before filing.
	Design the system to say "no source found" instead of guessing; show source links and confidence.	Rewards abstention over fabrication.
	Never use the same model to fact-check its own output.	The model confirmed its own fake cases.
Process	Require human review for high-stakes content (legal, medical, financial) with sign-off by a named person.	Keeps a human accountable for the final result.
	Keep an audit trail of prompts, outputs and which sources were verified and by whom.	Enables fast investigation and correction.
Organizational	Publish an AI usage policy: approved tools, what data may be entered, and what must be verified.	Sets clear expectations before problems occur.
	Train staff on how LLMs fail: hallucination, overconfidence and agreeing with leading questions.	Schwartz did not know the tool could invent content.
	Adopt a report-early norm: any suspected AI error is disclosed immediately and openly.	The cover-up, not the error, drew the harshest response.
Interface	Label AI-generated content and show warnings on tools used for professional research.	Reduces over-trust by non-expert users.

Suggested verification workflow

- Generate: the AI drafts an answer using retrieved, trusted documents only.
- Check: an automated verifier matches each citation to a real record and flags anything unmatched.
- Review: a qualified human confirms flagged items and spot-checks the rest.
- Record: sources and reviewer are logged before the output is released.

3. Responsible AI Summary

The *Mata v. Avianca* incident shows that fluent, confident AI output is not the same as accurate output. The tool behaved as language models do. The failure was using it in a high-stakes setting without grounding, verification or accountability, and then compounding the error by not being transparent once problems surfaced.

Principle	Lesson from the case
Accuracy and reliability	Ground outputs in verifiable sources and test systems for hallucination before deployment.
Transparency	Be clear about what is AI-generated and what has been independently checked.
Human accountability	A person owns the final result. "The AI said so" is not a defense.
Honesty about failure	Admit and fix errors quickly; concealment turned a mistake into sanctions for bad faith.

Key takeaway for AI engineers: design systems so that verifying output is easier than trusting it blindly. Retrieval grounding, citation checking and clear uncertainty signals turn a persuasive text generator into a tool professionals can safely rely on.

Sources

Mata v. Avianca, Inc., No. 22-cv-1461, 678 F. Supp. 3d 443 (S.D.N.Y. June 22, 2023) (Castel, J.).

Bloomberg Law, "Phony ChatGPT Brief Leads to \$5,000 Fine for NY Lawyers," June 22, 2023.

LawNext, "Court Imposes Sanctions on Lawyers Who Filed Bogus Cases After Relying on ChatGPT for Legal Research," June 2023.

CNN / ABC7, "Lawyer apologizes for fake court citations from ChatGPT," May 29, 2023.

Above the Law, "Lawyers Who Used ChatGPT To Fake Opinions Are In Real Trouble," June 2023.

Louisiana Law Blog, summary of the June 22, 2023 sanctions opinion.